

# Optimizing High-Performance Data Processing for Large-Scale Web Crawlers - News Straits Times (NST Online)

Prepared by:

VINESH A/L VIJAYA KUMAR A22EC0290

JOSEPH LAU YEO KAI A22EC0055

TIEW CHUAN SHEN A22EC0113

NUR FARAH ADIBAH BINTI IDRIS A22EC0245

# INTRODUCTION

The project is aimed to provide comparative analysis on the impact of different libraries to web scraping, data cleaning and analysis.

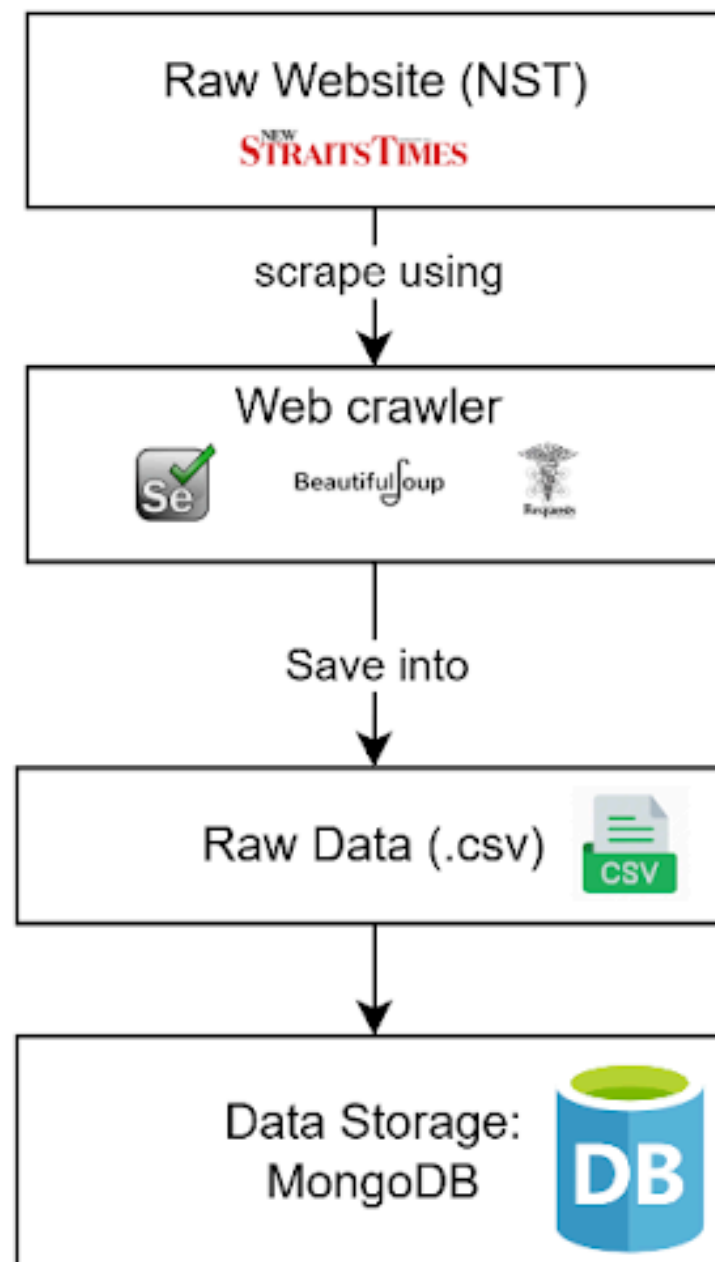
## Objectives

1. To develop a web crawler that is able to extract at least 100,000 records from a News Straits Times (NST) website.
2. To store extracted data in CSV format for further processing.
3. To clean and preprocess the raw dataset.
4. To evaluate performance before and after optimization using several performance metrics.

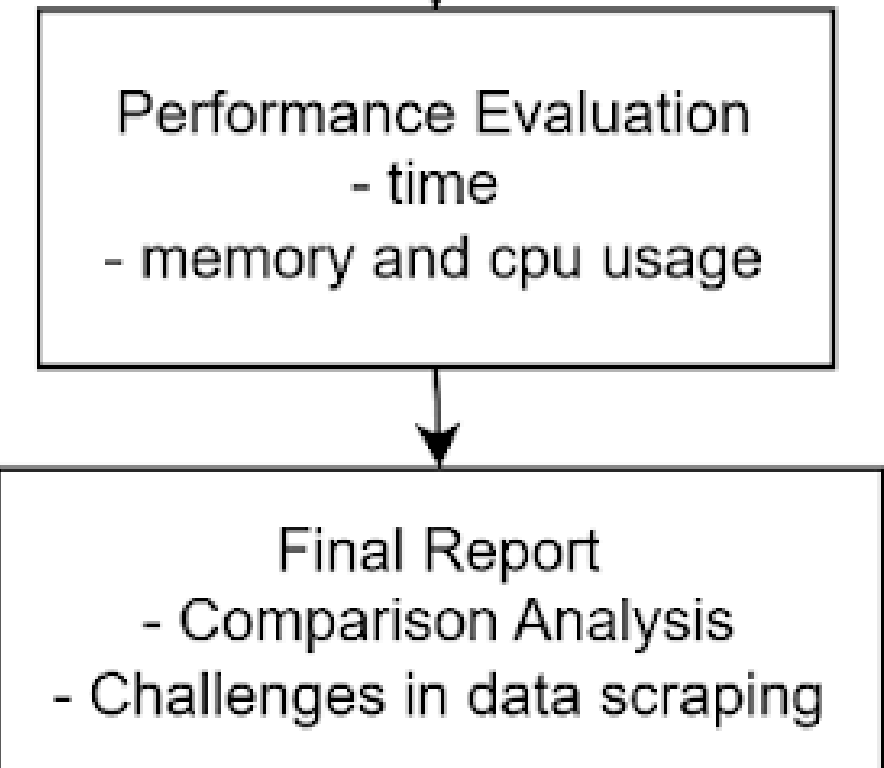
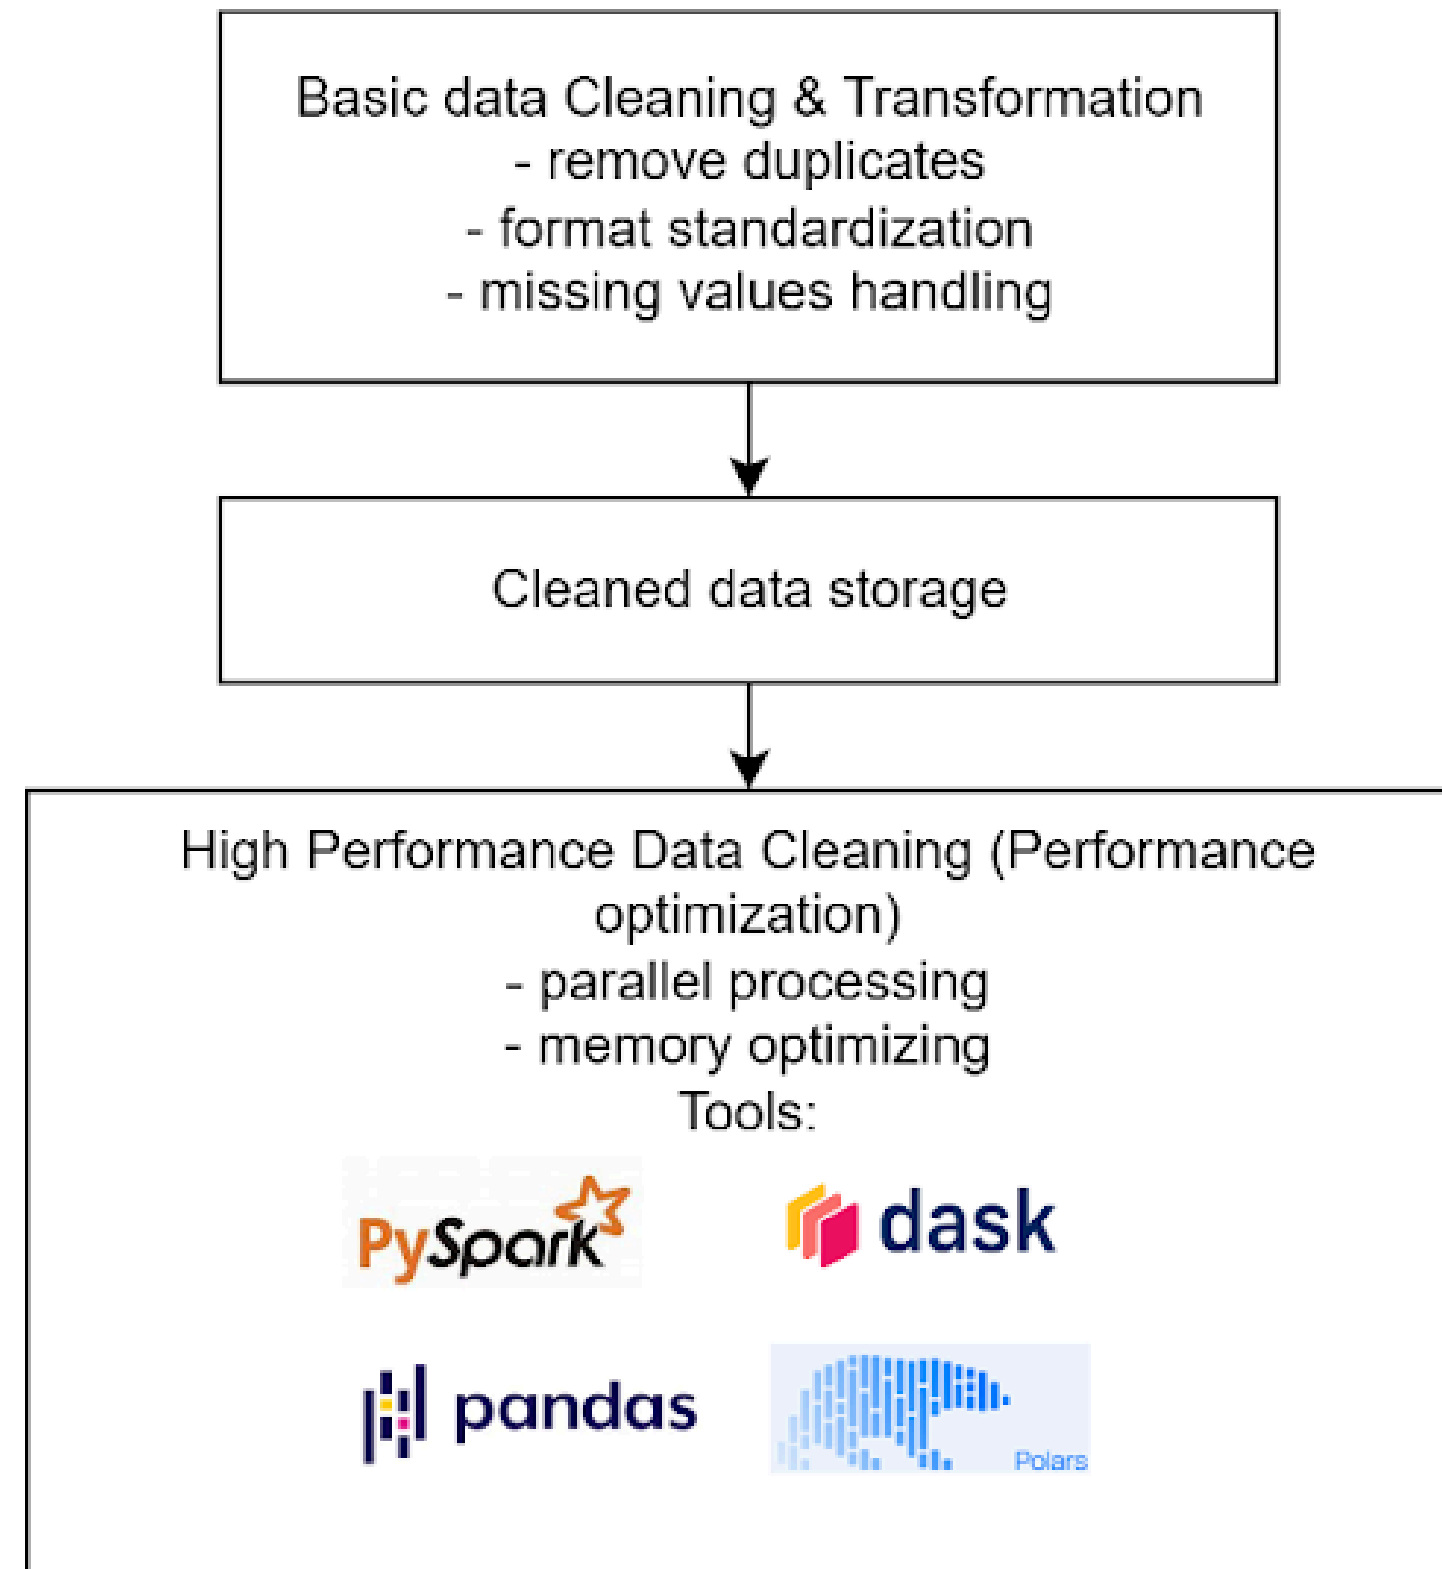
The logo for the New Straits Times newspaper. It features the word "NEW" in a small, black, serif font above the word "STRAITS" in a large, bold, red, serif font. To the right of "STRAITS" is the word "TIMES" in a large, bold, red, serif font. Above "TIMES" is the text "ESTABLISHED 1845" in a small, black, serif font.

# SYSTEM ARCHITECTURE

## Web Scraping



## Big Data Processing



# DATA COLLECTION

Data Field	Data Type	Description
Section	String	News topic(crime, politics, nation, health).
Publication date	Date	The date(including time) the article is published with format mm:dd:yyyy @ hh:mm
Headline	String	Title of the article.
Summary	String	Brief summary of the news.

- Data aimed to be be extracted from website
- Total rows of data extracted:127729

Crawling Method  
1.Pagination Handling  
2. Rate Limiting

## Ethical Consideration

- 1.Implemented delays between requests
- 2.Used headless browser configuration
- 3.Implemented error handling
- 4.Collected only publicly available information
- 5.Followed NST's robots.txt guidelines

# Data Processing

The data processing implementation is focused on data cleaning, transformation and storage using different libraries- Pandas, PySpark, Dask and Polars, for performance comparison and evaluation. Raw data is loaded from MongoDB for data cleaning process.


The cleaning and transformation methods

1. Null value
2. Duplicate value
3. Inconsistent capital/small number under the column-section
4. Wrong Date format




# OPTIMIZATION TECHNIQUE

**Polars**  
columnar data  
storage



**Dask**  
processing in parallel  
and lazily

**PySpark**  
runs on top of the  
Apache Spark  
framework(enables  
distributed computing)



**Vectorized  
Pandas**  
Pandas functions that  
operate on entire  
columns instead of  
using .apply()

The background features a dark blue field with a pattern of small squares in teal, pink, and orange, and thin white vertical lines of varying heights.

# ■ OPTIMIZATION CODE OVERVIEW

# Optimization: Polars

- **.str.split('@')**: Splits the string at the "@" symbol, creating a list.
- **.list.first()**: Extracts the first element of the list (the date part).
- **.str.strip\_chars()**: Removes leading/trailing whitespace.
- **.str.strptime(pl.Datetime, format='%b %d, %Y', strict=False)**: Parses the cleaned string into Polars' Datetime type based on the specified format. Polars' string operations are often faster than standard Python loops.

```
if 'Date' in df.columns:
    df = df.with_columns(
        df['Date']
        .str.split('@')
        .list.first()
        .str.strip_chars()
        .alias('Date')
    )

# Convert to datetime
df=df.with_columns(pl.col('Date').
str.strptime(pl.Datetime,
format='%b %d, %Y',
strict=False).alias('Date'))
```



# Optimization: Dask

- **.map\_partitions(lambda df: ...)**: Applies the function to each partition of the Dask DataFrame.
- **df['Date'].map(lambda x: x.split('@')[0].strip() if isinstance(x, str) else x)**: split and strip the date string.
- **pd.to\_datetime(df['Date'], errors='coerce')**: convert strings to datetime objects.

```
df_cleaned =  
df_cleaned.map_partitions(lambda df:  
df.assign(Date=df['Date'].map(lambda x:  
x.split('@')[0].strip() if isinstance(x,  
str) else x)))
```

```
df_cleaned =  
df_cleaned.map_partitions(lambda df:  
df.assign(Date=pd.to_datetime(df['Date'],  
errors='coerce')))
```

# Optimization: PySpark

- **regexp\_replace(col("Date"), "@.\*\$", ""):**  
Uses regular expressions to remove the "@" and everything after it.
- **regexp\_replace(col("Date"), "\s+", " "):**  
Normalizes multiple spaces into a single space.
- **trim(col("Date")):** Removes leading and trailing spaces.
- **to\_date(col("Date"), "MMM d, yyyy"):**  
Converts the cleaned string column to Spark's Date type using the specified format. Spark's operations are designed for large-scale distributed data processing.

```
# Remove time part
df_cleaned = df_cleaned.withColumn("Date",
    regexp_replace(col("Date"), "@.*$", ""))
```

```
# Normalize spaces
df_cleaned = df_cleaned.withColumn("Date",
    regexp_replace(col("Date"), "\s+", " "))
```

```
# Trim spaces
df_cleaned = df_cleaned.withColumn("Date",
    trim(col("Date")))
```

```
# Parse to date
df_cleaned = df_cleaned.withColumn("Date",
    to_date(col("Date"), "MMM d, yyyy"))
```

# Optimization: Vectorized Pandas

- **`df['Date'].str.split('@').str[0].str.strip()`:**  
These are vectorized string operations that apply the split, indexing, and stripping operations to the entire 'Date' Series efficiently without explicit loops.
- **`df['Date'].str.replace(r'\s+', ' ', regex=True)`:** Another vectorized string replacement using regular expressions for efficient space normalization.
- **`pd.to_datetime(df['Date'], errors='coerce')`:** converting entire Series of strings to datetime objects.

```
if 'Date' in df.columns:
```

```
    df['Date'] =  
    df['Date'].str.split('@').str[0].str.strip()  
    df['Date'] = df['Date'].str.replace(r'\s+',  
    ' ', regex=True)
```

```
    df['Date'] = pd.to_datetime(df['Date'],  
    errors='coerce')
```

The background is a dark blue field decorated with various geometric elements. There are several thin white vertical lines of varying lengths. Scattered throughout are squares of different colors: teal, orange, and light blue. Some squares are solid, while others are outlined. The text 'PERFORMANCE EVALUATION' is centered in the middle of the image.


# PERFORMANCE EVALUATION

# Web Scrapping

Operation	Aspects	Comparisons				
		Dask	Polars	Pyspark	Pandas	Vectorized Pandas
Dataset Loading and Display	Code Execution Time (s)	0.5574	0.13728	0.1715	1.45037	0.81039
	Peak Memory Usage (MB)	9.566	0.3828	0.0	0.0	0.0
	Throughput (rows/s)	217847.737	884573.6415	707961.62	83728.7675	149849.72

***Table 2: Comparison between Data Processing and Cleaning Techniques***

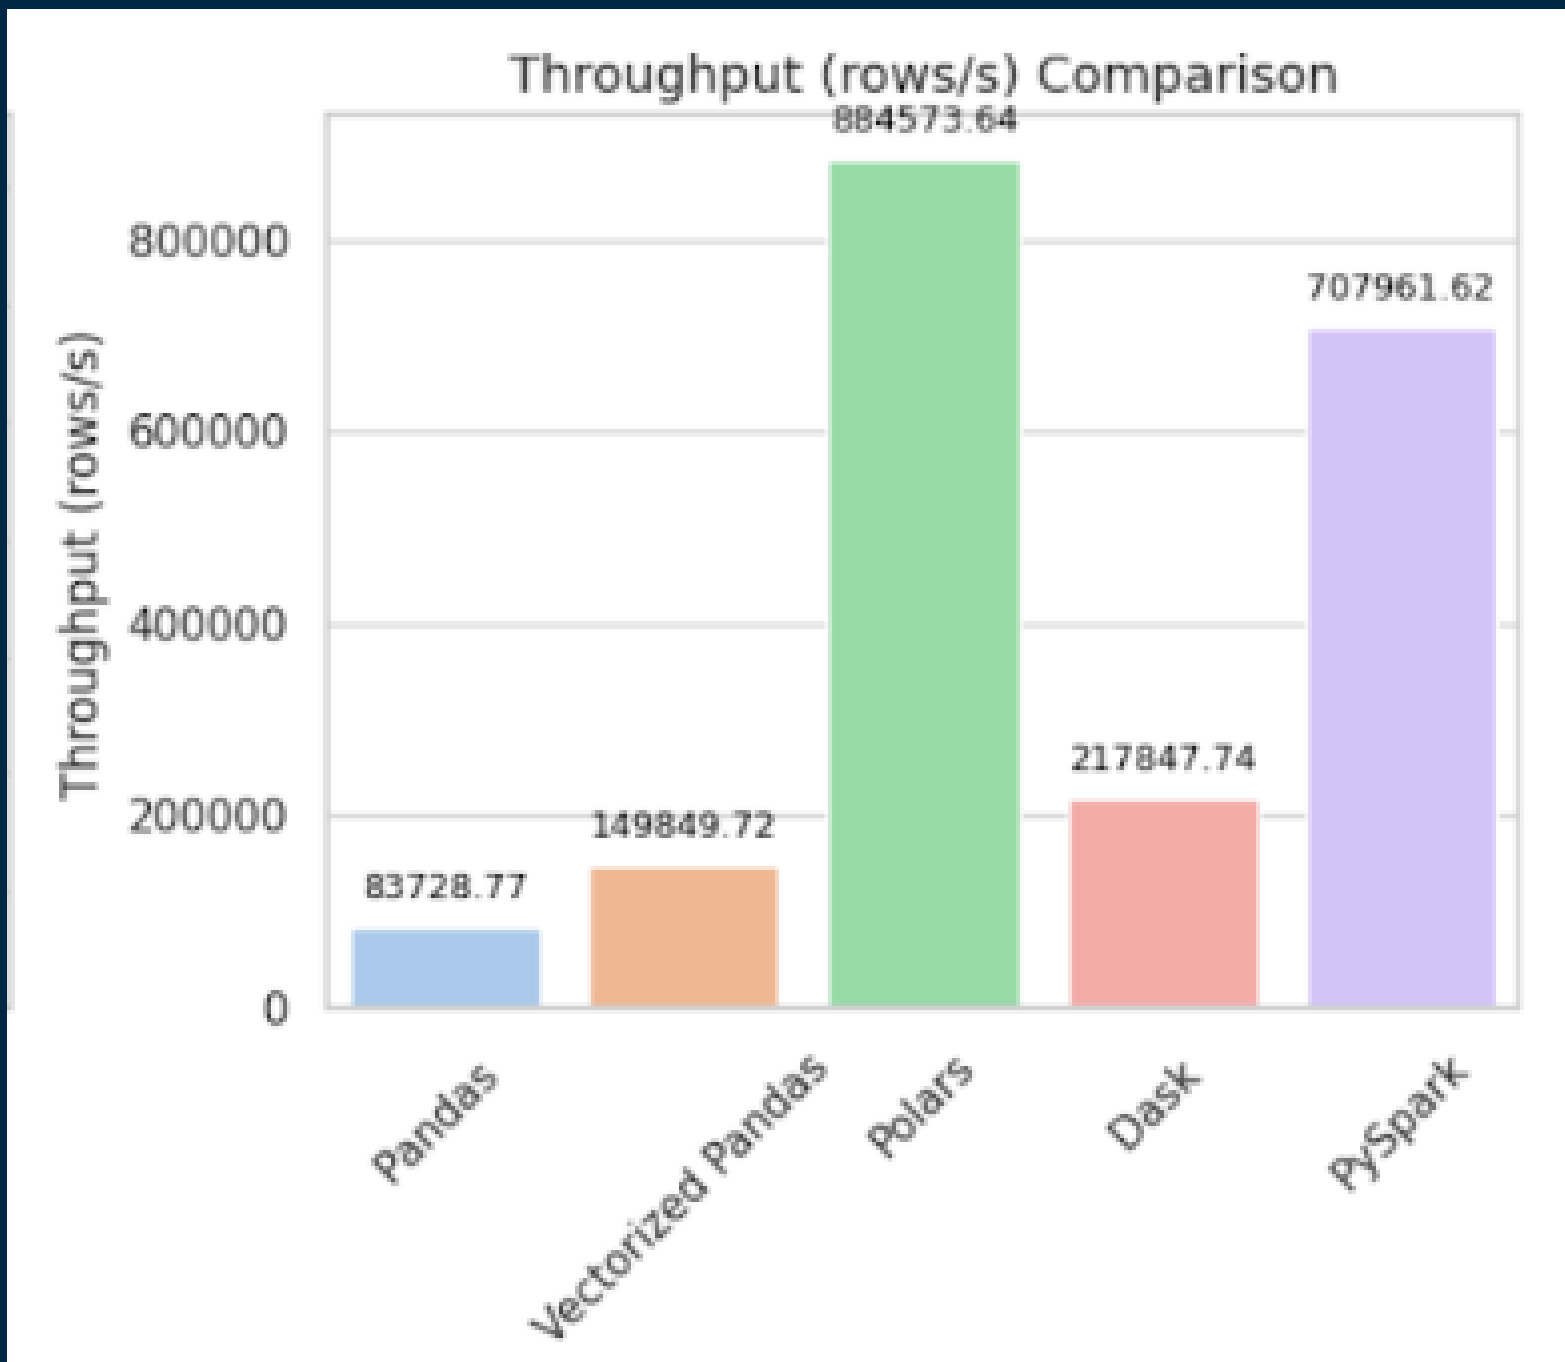
# Web Scrapping-Time(s)



## Performance Comparison:

- Vectorized Pandas (1.45s)
- Pandas (0.81s)
- Polars (0.14s)
- Dask (0.56s)
- PySpark (0.17s)

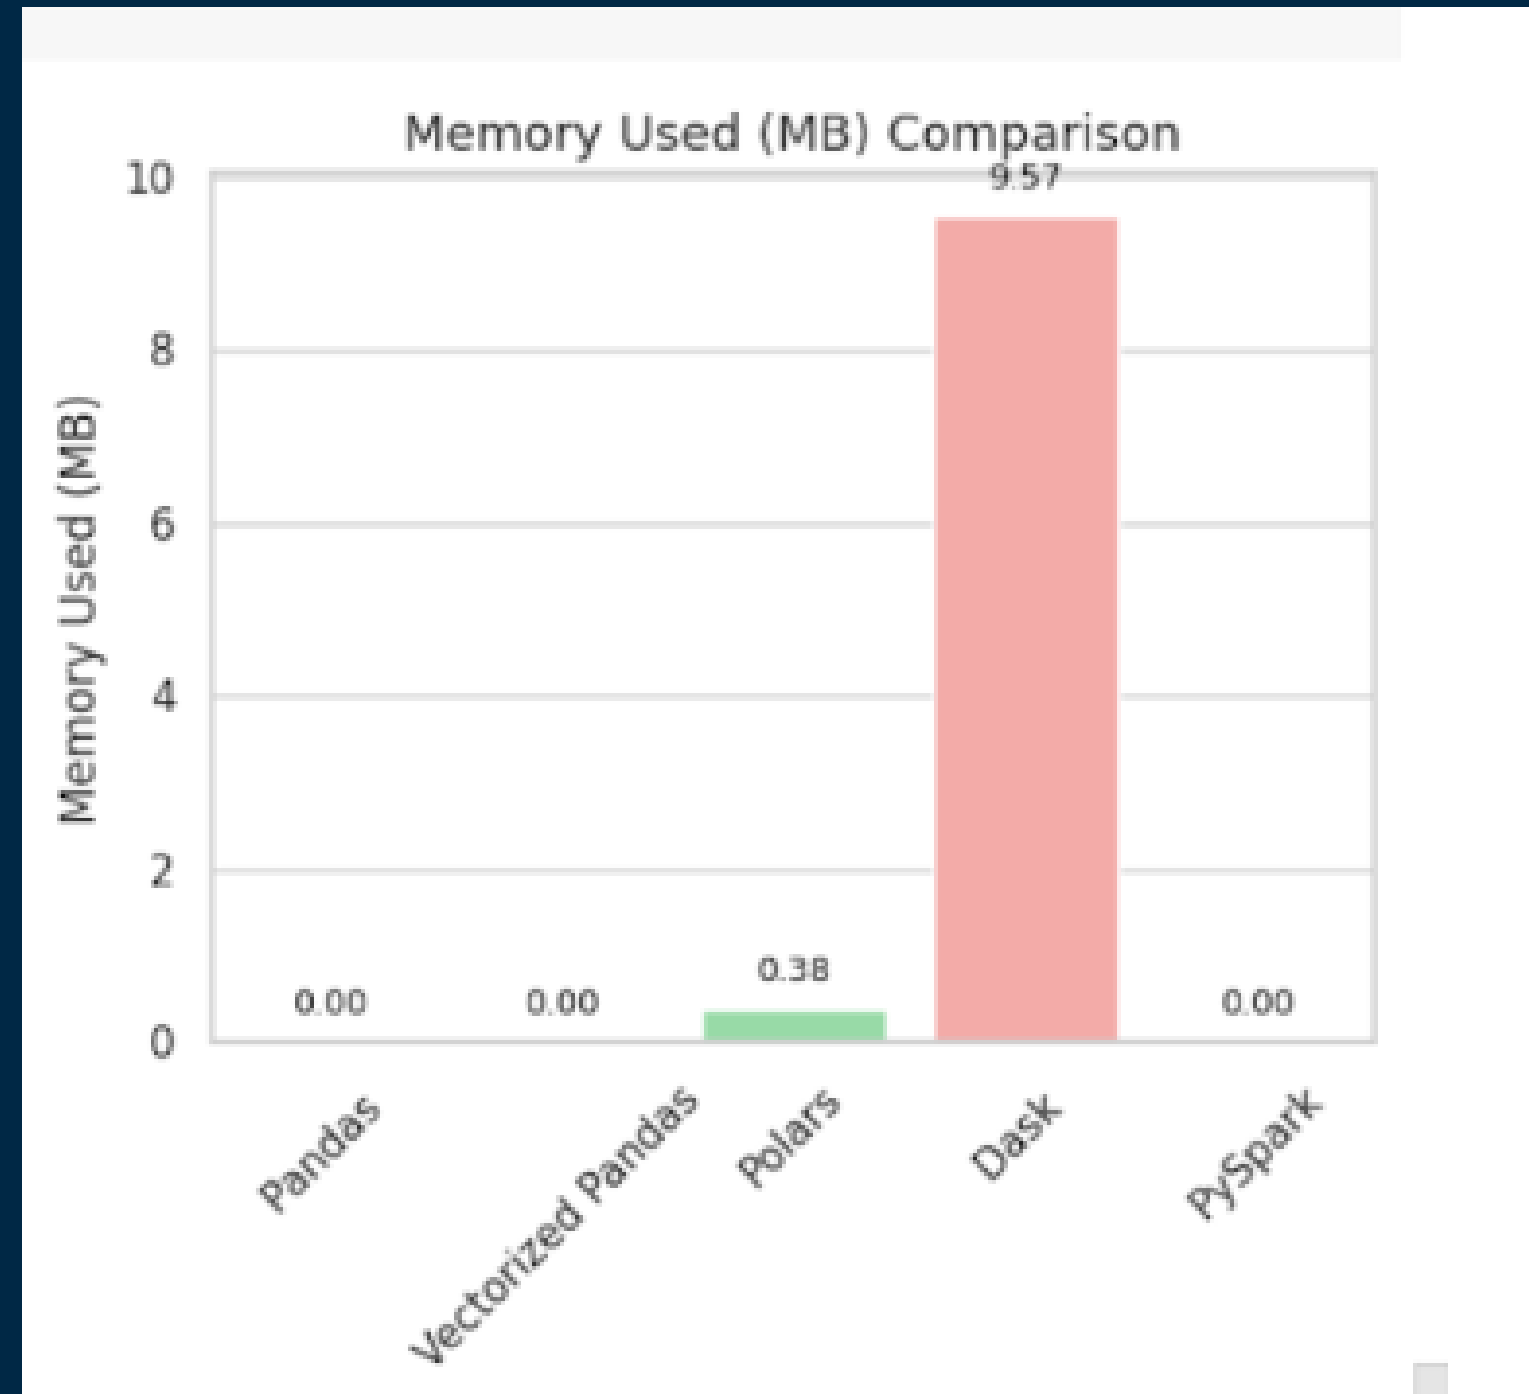
# Web Scrapping-Throughput(rows/s)



## Throughput Comparison:

- Polars (884573.94 rows/s)
- PySpark (707961.62 rows/s)
- Dask (211847.74 rows/s)
- Vectorized Pandas (140520.72 rows/s)
- Pandas (8528.77 rows/s)

# Web Scrapping-Memory(MB) usage



## Memory Usage Comparison:

- Dask consumes the most memory (9.57 MB)
- Polars (0.38 MB)
- Pandas, Vectorized Pandas, and PySpark use (0.00 MB)



The background is a dark blue field decorated with a pattern of small squares and thin white lines. The squares are in various colors: teal, light blue, orange, and pink. Some squares are solid, while others are just outlines. The lines are thin and white, some horizontal and some vertical, creating a grid-like but irregular pattern.

# ■ CHALLENGES & ■ LIMITATIONS


# Web Crawler for NST Website – Overview & Challenges

- Developed to scrape NST articles; initially used Scrapy, BeautifulSoup, Selenium.
- Fully shifted to **Selenium** due to dynamic content.
- **Key Challenges:**
  - Selenium reliable but **resource-heavy** and **slow**.
- **Ethical scraping:** single-user mode with delays – low scalability.
- **Data cleaning issue:** inconsistent 'Date' formats.
- **Pandas** too slow for large datasets (~100k+ rows).

CPU ?



Memory ?



# Performance, Bottlenecks & Improvements

- **Optimization Testing:**

**Polars:** fastest for large datasets.

**Dask:** high memory usage.

**PySpark:** slow to start.

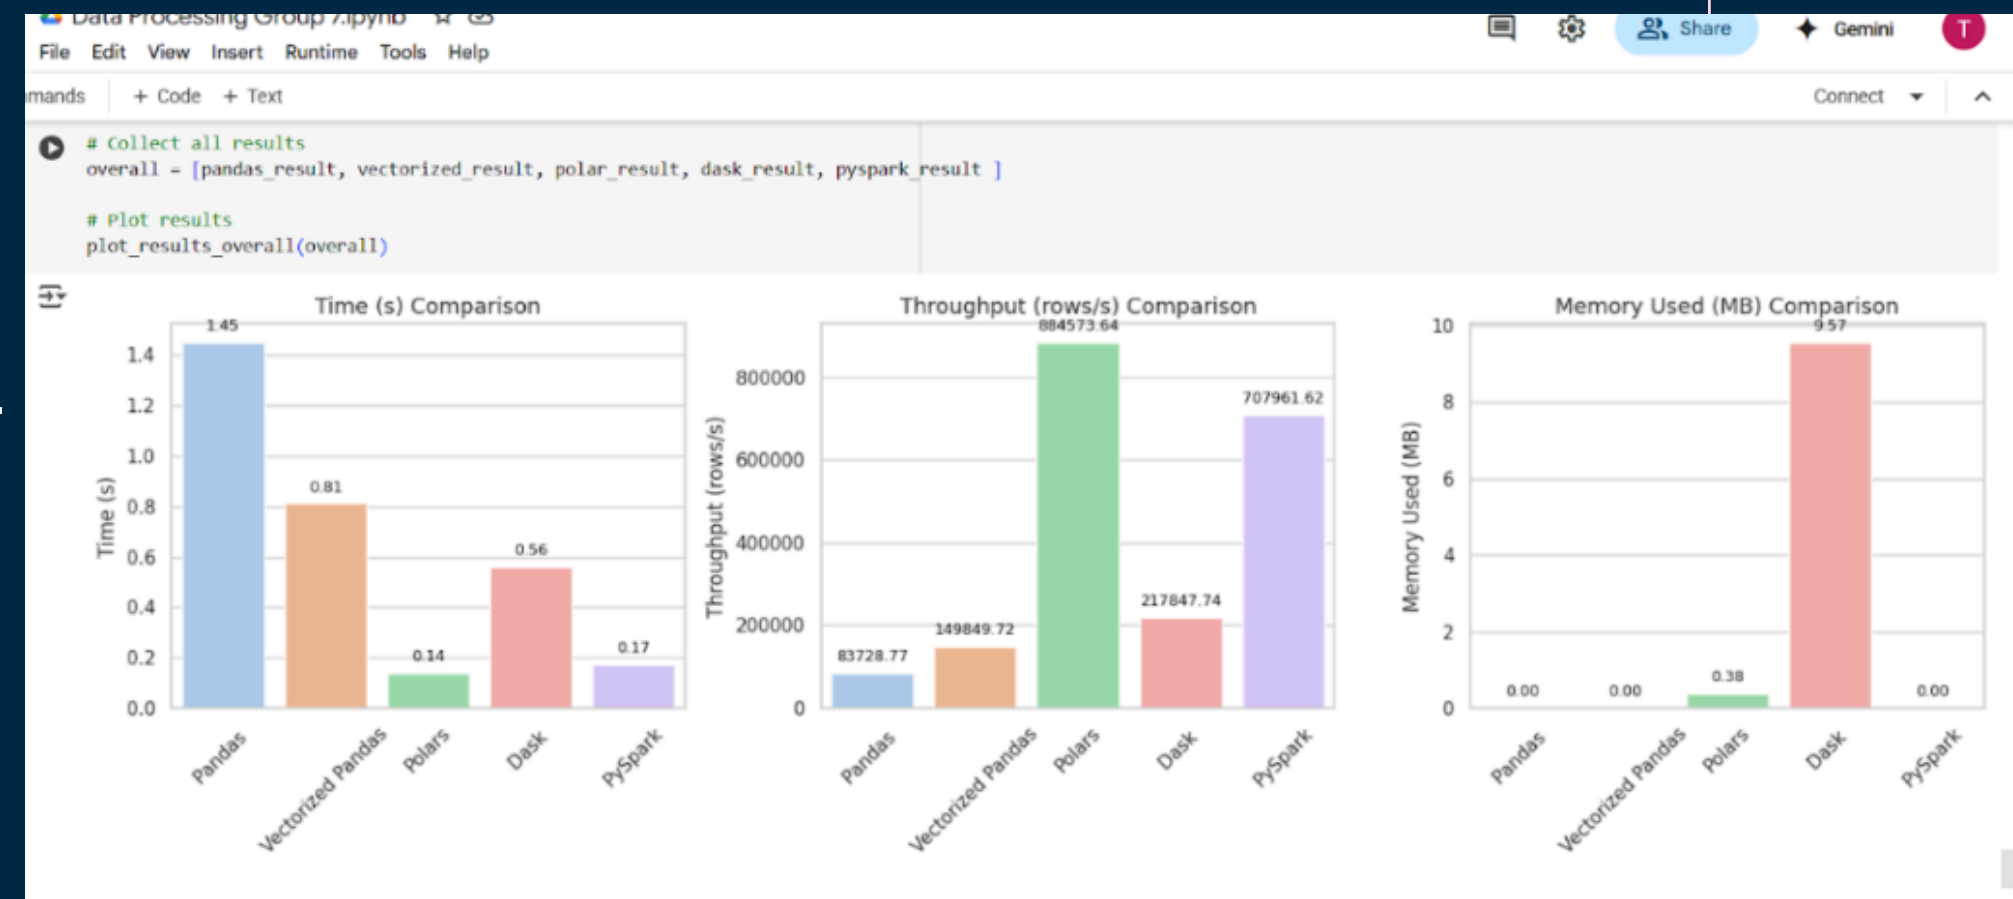
**Vectorized Pandas:** average speed.

- **Limitations:**

Hardcoded to NST's structure, not reusable.  
No user interface or hardware-aware logic.

- **Future Improvements:**

Flexible scraping engine.  
Modular, site-agnostic pipelines.  
GUI for non-technical users



The background is a dark blue field decorated with a pattern of small squares and thin white lines. The squares are in various colors: teal, orange, pink, and light blue. Some squares are solid, while others are outlined. The lines are thin and white, some running vertically and others horizontally, creating a grid-like structure. The text is centered in the middle of the image.

# CONCLUSION & FUTURE WORK

# Polars Library



Time Comparison

- 0.14 seconds

Throughput

- 884573.64 rows per second

Memory Usage

- 0.38 MB

If compared to pandas, Polars can achieve more than 30x performance gains.

01

# Future Work



- More complex data transformation
- multiple runs to account for variability and more reliable averages
- larger dataset to stress-test the methods
- investigation into why Dask's memory usage was higher

02